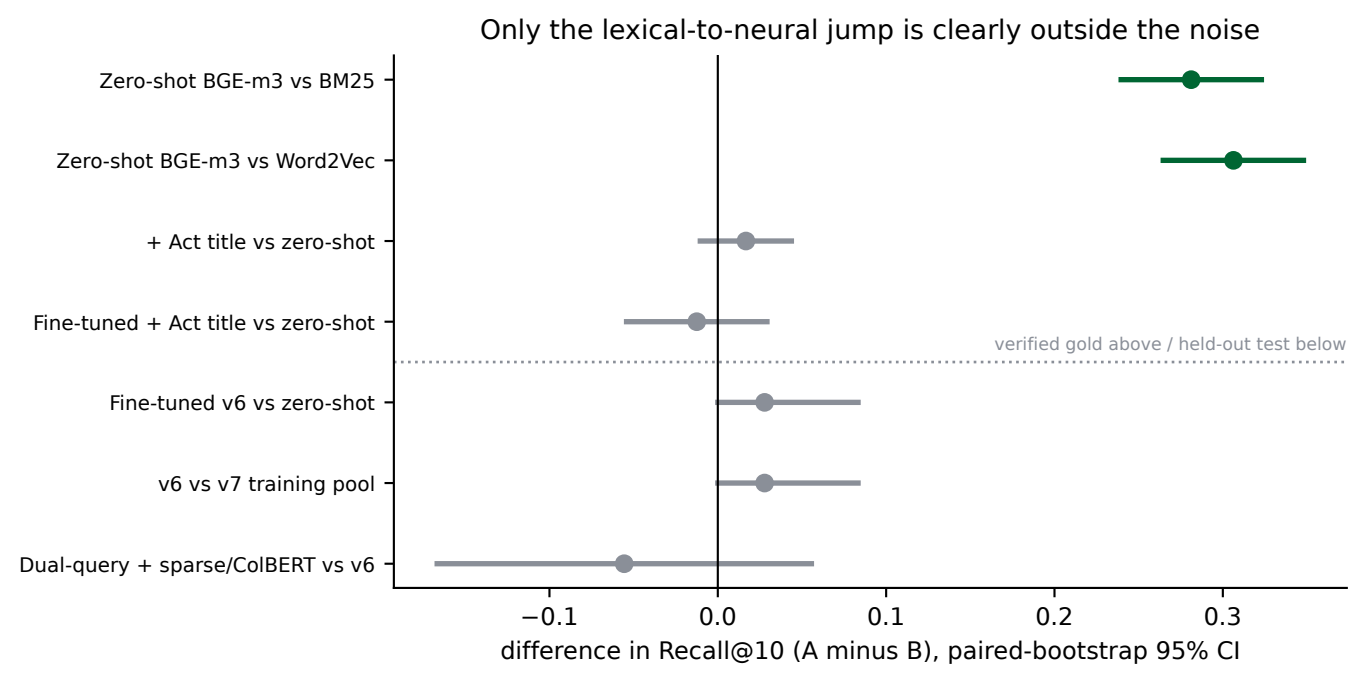
Only the lexical-to-neural jump is clearly outside the noise
Zero-shot BGE-m3 vs BM25
Zero-shot BGE-m3 vs Word2Vec
+ Act title vs zero-shot
Fine-tuned + Act title vs zero-shot
verified gold above / held-out test below
Fine-tuned v6 vs zero-shot
v6 vs v7 training pool
Dual-query + sparse/ColBERT vs v6
−0.1
0.0
0.1
0.2
0.3
difference in Recall@10 (A minus B), paired-bootstrap 95% CI